

A paid application standardises what large language models read about insulin pump settings and leaves what they decide to chance: reproducibility of an automated insulin delivery settings review across eleven models

Tim Street, MEng BEng (Hons), ORCID 0009-0008-4417-6581. Diabettech Ltd, London, United Kingdom. tim@diabettech.com

Preprint, not peer reviewed. Posted to SSRN and diabettech.com in August 2026.

Abbreviations: AID, automated insulin delivery; CGM, continuous glucose monitoring; DIA, duration of insulin action; IOB, insulin on board; ISF, insulin sensitivity factor; JSON, JavaScript Object Notation; LLM, large language model.

Keywords: AndroidAPS; artificial intelligence; automated insulin delivery; large language models; reproducibility; type 1 diabetes

Abstract

Background. A paid Android application prepares a user's AndroidAPS data and four prompts for any large language model, checks the returned file and presents it as a settings review.

Method. The application's package and prompts were replayed verbatim against eleven models, 50 conversations per model (20 for two), on one adult's 14-day record. Responses were judged by a validator verified against the application's behaviour. Outcomes were whether quoted figures were in the record, agreement on 23 settings within and between models, and what the check accepted, with Wilson 95 per cent intervals.

Results. The models quoted the record accurately: 96.8 to 100 per cent of cited figures were in it. Three models lowered the maximum insulin on board limit in 73 to 98 per cent of accepted conversations, proposing 5 to 18 distinct value pairs; four left it alone throughout. One model proposed sensitivity factors from 30 to 180 mg/dl per unit against a current 56. Acceptance at first paste ranged from 34 per cent (22.4 to 47.8) to 100 per cent, and up to 17 of 20 accepted responses broke rules the prompts state but the check does not enforce.

Conclusions. A sensitivity factor three times the current value and a meal step marked completed on a record with no meals both passed the application's check. Which settings a user is told to change depended more on the model opened than on the data, and the user is given no means of telling which answer they have.

Introduction

A person who runs an open-source automated insulin delivery (AID) system such as AndroidAPS sets their own basal rates, carbohydrate ratios, insulin sensitivity factors (ISF), duration of insulin action (DIA) and glucose targets, and the limits that bound what the algorithm may

deliver [1]. Reviewing those settings against a fortnight of glucose data is skilled work, and some users now ask general-purpose large language models (LLMs) to do it between clinic appointments [2, 3]. By August 2026 an Android application had made the practice a product [4]. It builds a JavaScript Object Notation (JSON) package from the user's Nightscout instance and AndroidAPS profile, with a decision inventory of 23 settings, and supplies four prompts. The user opens a new conversation in whichever LLM they prefer, attaches the package with the first prompt, sends the other three in turn and pastes the final JSON response back into the application. The application checks it, returns the user to the conversation with a repair prompt if it fails, and otherwise renders a primary profile recommendation, a decision for every setting, a meal strategy and a list of recommendations; a subscription costs £4.49 a month, and no model is recommended.

The product rests on the promise that a model given a well-structured record will derive a review from it, and on the assumption that checking the review's form is enough to stand behind its content. General-purpose models answer diabetes questions fluently and sometimes wrongly [5], and their carbohydrate estimates from photographs vary from query to query [6, 7, 8]. Asked to propose pump settings from a week of raw continuous glucose monitoring (CGM) and pump data, five models anchored on textbook values unless told the user's current settings, and 30 to 70 per cent of the glucose readings they cited were not in the data; the authors concluded that a deployment would need two safeguards, the user's settings as context and programmatic checking of what the model cites [9]. The Diabetes Technology Network UK has advised that any suggested change to a pump setting must be checked for plausibility before it is acted on [10]. The application supplies the first safeguard, since the package carries the current profile, and its check is a form of the second; what happens to the review once both are in place has not been measured.

This study replicated the application's workflow against eleven models, on one real record, and asked three questions in order: whether the models read the record correctly, whether they agreed about what to change, and what the application's check made of the difference. The product's regulatory position is taken up in a companion commentary.

Methods

Materials

The application studied was TuneMyBG, version 1.0.5 (Adam Kowalczyk, Lubliniec, Poland; distributed through Google Play) [4]; the package and the four prompts were copied verbatim from it. The application's source is not available, so its repair prompt and acceptance behaviour were established by pasting 25 real model outputs into the result screen and recording what it did. The application's synthetic demonstration package served as a control. The study package was generated by the application from the author's own Nightscout instance and AndroidAPS profile: 14 days of CGM data, mean glucose 121.6 mg/dl, time in range 86.0 per cent, time below 70 mg/dl (3.9 mmol/L [11]) 5.1 per cent, 275.95 units of bolus and SMB insulin in its treatment records, and no carbohydrate entries, since the author does not log meals. The profile had a single ISF of 56 mg/dl per unit, overnight targets of 90 mg/dl, a DIA of 10 hours, a maximum basal of 12 U/h and a maximum insulin on board (IOB) of 25 U. Time below 70 mg/dl peaks at

18.5 per cent after midnight and has fallen to low single figures by 04:00 with the same basal segment running, which under a 10-hour DIA is more consistent with the tail of evening insulin than with the overnight basal or target; that reading is the author's and is offered as interpretation.

One record was used by design. The question is whether the same input yields the same review, and holding the input fixed while the conversation is repeated attributes all observed variation to the model and the application. A study of whether the advice is correct would need many records and a reference standard, and is a different study. The data are the author's own; he gave written consent to their use and publication, and no other person's data were involved, so no institutional review board approval was sought. No change was made to any treatment.

Replication of the workflow

Each conversation was fresh, with the package attached to the first prompt and the remaining prompts sent in turn. No sampling parameters were set, so each provider's default applied. Eleven models were run on 21 and 22 August 2026: Gemini 3.6 Flash, Gemini 3.1 Pro Preview and Gemini 3.5 Flash-Lite (Google, Mountain View, California, USA); GPT-5.6-sol and GPT-5.4-mini at the medium reasoning setting (OpenAI, San Francisco, California, USA); Claude Opus 5, Claude Sonnet 5 and Claude Haiku 4.5 (Anthropic, San Francisco, California, USA); DeepSeek V4 Pro (DeepSeek, Hangzhou, China); Grok 4.6 (xAI, Palo Alto, California, USA); and Llama 4 Maverick (Meta, Menlo Park, California, USA), the last three through OpenRouter (New York, New York, USA). Fifty conversations were run per model, except Opus and Sonnet at 20 because of cost, and the served model identifier was recorded for every turn (Supplementary Table S1). Conversations ran through the providers' programming interfaces, so the results describe the models rather than any chat product.

The application's check and the repair loop

The application issues a repair prompt for three conditions: decision rows missing from or absent in the inventory, a current value not copied exactly, and a primary recommendation focus outside the five allowed words. It stops without a prompt when the paste is not valid JSON or a required string field is absent. It accepts everything else, including invalid enum words, missing sections and extra keys. The validator built from these observations reproduces the application's verdict on all 25 informative test pastes (Supplementary Table S2), and it, rather than the application, judged the study's responses. A repair prompt was sent once in the same conversation when needed, and each conversation was classified as accepted at first paste, accepted after repair, rejected with no prompt offered, or still rejected after repair; rules stated in the prompts but not enforced were recorded separately. Content analyses cover accepted conversations only, since the application holds nothing from a rejected one.

Outcome measures

Whether the models read the record was measured by checking every figure with a unit quoted in an accepted response against the numbers in the package, and by recording whether each response recognised the absence of carbohydrate data and how it scored the meal strategy step. Whether the models agreed was measured, for each of the 23 settings, by the proportion of accepted conversations proposing a change and the values proposed, and by a stability index, the

mean proportion of conversations agreeing with the modal decision per setting. What the check made of it was measured by the acceptance classes and the proportion of accepted responses breaking an unenforced rule. Proportions carry Wilson 95 per cent intervals with the conversation as the unit. Within-model results are solid; the between-model comparison is provisional, since the models are versioned services; the reasons offered for the differences are interpretation. All analyses were specified before the runs except the carbohydrate-handling examination, the citation check, the per-decision check of whether a change's rationale cited a figure from the record, and the reading of the overnight rationales for whether they attributed the lows to earlier insulin or to the setting itself, and the examination of how responses read the package's insulin total, all added after inspection of the outputs.

Results

What the models read

On the control package all 50 Gemini 3.6 Flash conversations kept every setting unchanged; the remainder concerns the study record. The models quoted the record accurately (post hoc analysis): of the figures with a unit in accepted responses, 96.8 per cent (Opus) to 100 per cent (GPT-5.4-mini, Llama) were present in the package. Where a response changed a setting, the rationale cited at least one figure from the record in 85 to 100 per cent of changes for eight models, in 69 per cent for Gemini 3.1 Pro and in 5 per cent for Llama (Supplementary Table S3); no rationale quoted a figure that was absent from the record. Each model recognised that no carbohydrates had been logged in at least 70 per cent of its conversations (post hoc analysis). Only Llama changed a carbohydrate ratio (2 of 44 conversations); the others said the ratios could not be tested without meal data. The three models that acted on the safety limits cited the same arithmetic, a 25 U ceiling against 19.7 U of insulin a day, and both branches of the fork below cited the same figure, 18.5 per cent of readings below 70 mg/dl at 01:00. The 19.7 U is inherited: the package's insulin total sums only the bolus and SMB records, it carries no basal delivery, and its scheduled basal adds 16.2 U a day (post hoc analysis; Supplementary Table S4). Flash took the bolus-only sum as the daily total in 49 of 50 conversations and lowered the limit in 48; GPT-5.4-mini, Haiku and Opus noted the omission in 67 to 95 per cent of theirs; Opus reconstructed about 36 U a day in all 11.

What the models concluded

Gemini 3.6 Flash, Gemini 3.1 Pro and Opus lowered the IOB limit in 98, 90 and 73 per cent of accepted conversations; four models left it alone throughout; DeepSeek and Llama lowered it in 44 and 39 per cent. Four models changed nothing in most conversations, GPT-5.4-mini and Sonnet among them.

Within a model the direction of a change was usually consistent and the number was not (Figures 2 and 3). Gemini 3.6 Flash lowered the IOB limit in 49 of 50 conversations (98 per cent, 89.5 to 99.6) and the basal limit in 48, proposing 18 distinct pairs of values. Gemini 3.1 Pro produced 17 distinct pairs and Opus, in 8 of 11 conversations (73 per cent, 43.4 to 90.3), five. Stability indices ranged from 0.915 (DeepSeek) and 0.928 (Haiku) to 0.998 (GPT-5.4-mini and Sonnet); an index of 1 would mean the same decision on all settings in all conversations; for the least stable model roughly one decision in twelve departed from the modal answer.

The primary recommendation also forked within a model: Gemini 3.6 Flash raised the two overnight targets from 90 to 100 mg/dl in 39 of 50 conversations (78 per cent, 64.8 to 87.2) and headed the result with “target”; in the other 11 it left the targets alone, cut the midnight basal from 0.6 to 0.5 U/h and headed the result with “basal”. Llama changed the midnight basal in 23 of 44 accepted conversations, raising it to 0.65 U/h in 16 and lowering it to 0.55 U/h in seven. DeepSeek changed the ISF in 20 of 48 (42 per cent, 28.8 to 55.7), proposing 14 distinct values from 30 to 180 mg/dl per unit against a current 56.

The overnight lows themselves were read three ways (Table 3; post hoc analysis). Opus and Sonnet traced them to earlier insulin in 11 of 11 and 18 of 19 conversations, and neither changed the overnight target or basal in any conversation. Gemini 3.1 Pro traced the lows to “massive late insulin stacking from evening UAMs” in 40 of 50 conversations and still raised the target in 45, 36 of them while naming the stacking. Llama offered no causal attribution in 18 of its 23 basal changes.

The meal strategy step, which had no meal data to work from, was marked completed by Flash-Lite, Llama and Haiku in 39 of 39, 43 of 44 and 20 of 20 conversations, and partial or blocked by GPT-5.6-sol, Opus and Sonnet in all but one.

What the check caught

Table 1 gives the acceptance classes with intervals. Opus produced output that was not valid JSON in nine of 20 conversations (45 per cent, 25.8 to 65.8), and Flash-Lite in 10 of 50 (20 per cent, 11.2 to 33.0); the application offers no way forward from that state. Llama recovered through the repair prompt in 25 of 50 conversations, whereas Haiku was still rejected after the repair prompt in 27 of 50 (54 per cent, 40.4 to 67.0).

Nothing in the preceding section was caught: the check accepted all 40 safety-limit pairs from the three models that lowered the limits, both branches of the Flash fork, both directions of the Llama basal change, all 14 DeepSeek sensitivity factors and each meal step marked completed on a record with no meals. Among accepted responses the proportion breaking a rule the prompts state but the check does not enforce ranged from none (GPT-5.6-sol) to 25 of 48 (GPT-5.4-mini) and 17 of 20 (Haiku), the commonest breaches being confidence words outside the allowed set, a missing issues section and extra keys, which reach the result screen.

Discussion

On this record the review a user received depended more on which model they opened than on the data, and the application could not tell the difference. The models did not agree about what followed from the record: four models left the limits alone, three lowered them, one proposed sensitivity factors that ranged over a factor of six, and one raised a basal rate that another lowered. Within a single model the direction held and the number did not. The overnight lows show the difference at its clearest: the same figure led two models to leave the overnight settings alone because earlier insulin explained it, two to raise the target as a cushion against it, and one to change the basal without offering a cause; only the first acts on what the record shows. Each answer was internally coherent and presented in the same form, so nothing on the screen tells the user which they have been given.

This is a different failure from the two described for models given raw five-minute data, where settings anchored on training priors and the cited readings were unreliable [9]. Here the package carried the current profile and the check enforced the copying of current values, the two safeguards that paper proposed. The misquotation did not recur, and for eight of the eleven models nearly every change was argued from a figure in the record, for Gemini 3.1 Pro two thirds were; Llama changed settings without citing one, and DeepSeek's shortening of the DIA has the shape of the textbook anchoring seen before. The comparison is between two studies with different records, prompts and models, so it shows the earlier failures largely absent here, not that the package removed them. The package also fixes what the models do not see: its insulin total omits basal delivery, and the model that cut the safety limits most often was the one that mistook the bolus-only sum for the whole. The models that read the total correctly did so from training, since knowing that an AID system delivers basal insulin is not in the package; the priors the earlier study treated as anchoring are also a model's means of questioning the data, and whether a draw applies them is itself sampled. What the models argued from the record was not stable: with good evidence they generate several plausible conclusions and the sampling decides which is shown. A structured package appears to fix what the model reads and to leave untouched what it decides, and the product sells the decision.

The application's check is a check of form, enforcing the identity of the decision table, the copying of current values and one enumerated field. It does not inspect a proposed value, a status, a confidence word or the content of a recommendation, so a meal step marked completed on a record with no meals appears as the model wrote it, and an ISF of 180 mg/dl per unit passes as readily as one of 56. The prompts contain rules against most of this, but they are instructions to the model rather than checks, and the cheaper models broke them in a quarter to four fifths of accepted responses.

None of this can be cured by a temperature setting, since the chat interfaces expose none and provider-side non-determinism persists at temperature zero [12]. A design that took the variation seriously would run the analysis several times and show the spread, and would validate and pin one model rather than accept any. The advice that every suggested change be checked for plausibility before it is acted on [10] assumes a user who can recognise which of several plausible answers they have been handed, and the application's screen gives them no means of doing so.

A profile with different problems might be treated differently. The validator reproduces the application's behaviour on 25 informative tests but could differ on cases not tested. The study measured consistency, not correctness: there is no reference answer for these settings and no counterfactual glucose trajectory for any proposed change, so what is reported is what the models produced, not what would have happened had a user acted on it. The safety concern is therefore epistemic rather than a claim that any particular answer was mistaken.

Conclusions

Eleven models quoted one glucose record accurately and returned reviews that differed in substance between models and in detail within them, and the application that sells the review accepted all of them. A single review from such a product is one sample from a distribution whose shape the user is not shown. A developer who sells one answer should first measure the

spread and disclose it, and a user who receives one should treat it as a hypothesis rather than a finding.

Funding

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors. Model usage for the study, \$158.42 plus £19.73, was paid for by the author.

Acknowledgments

Claude (Anthropic, San Francisco, California, USA), a large language model assistant, was used under the author's direction to write the harness and analysis code, to run the analyses, to prepare the figures and to draft and revise the text. The author specified the study, made every decision of design and interpretation, checked each reported number against the stored outputs, and is responsible for the content.

Disclosures

The author runs Diabettech Ltd, is a committee member of the Diabetes Technology Network UK and wrote its statement on large language models, and is the author of the two SSRN preprints cited. He has no relationship with the developer of the application studied and paid for its subscription. Three of the eleven models studied are Anthropic's, as is the assistant used to produce this work; Anthropic had no role in the study, was not informed of it, and all model usage was paid for at list price.

Data accessibility

The prompts, the harness, the validator, the analysis code, every transcript and model output, and the package with identifiers removed are published in the study repository. The package is the author's own health data and is released with his written consent.

References

1. AndroidAPS documentation. <https://androidaps.readthedocs.io>. Accessed 21 August 2026.
2. Street T. In conversation with: factual questions for large language models. Diabettech. 21 June 2025. <https://www.diabettech.com/ai/in-conversation-with-factual-questions-for-large-language-models/>. Accessed 22 August 2026.
3. Street T. ChatGPT "Health" conversations and a Boost AndroidAPS case study. Diabettech. 8 January 2026. Accessed 22 August 2026.
4. Kowalczyk A. TuneMyBG. Google Play. Version 1.0.5, updated 20 August 2026. <https://play.google.com/store/apps/details?id=com.adamkowalczyk.tunemybg>. Accessed 22 August 2026.
5. Sng GLY, Tung JYM, Lim DYZ, Bee YM. Potential and pitfalls of ChatGPT and natural-language artificial intelligence models for diabetes education. *Diabetes Care*. 2023;46:e103-5.

6. Goncalves S, Coelho C, Pretre L, et al. Chat, Gemini and Claude at the dinner table: assessing general-purpose AI tools for carbohydrate counting in type 1 diabetes. *Diabetes Res Clin Pract.* 2025;113031. <https://doi.org/10.1016/j.diabres.2025.113031>
7. Joubert M, Dreves B, Arnould T, et al. Comparative accuracy of smartphone apps and a generative AI tool for carbohydrate counting: an independent bicentric study. *Diabetes Obes Metab.* 2026. <https://doi.org/10.1111/dom.70396>
8. Street T. Reproducibility and accuracy of large language model vision APIs for carbohydrate estimation from food photographs: a four-model batch comparison with implications for automated insulin dosing. SSRN preprint. 2026. <https://doi.org/10.2139/ssrn.6577780>
9. Street T. Frontier LLMs require explicit clinical context to avoid training-data anchoring for insulin pump settings: a pre-registered exploratory study. SSRN preprint. 2026. <https://doi.org/10.2139/ssrn.6619638>. Pre-registration: OSF, 10.17605/OSF.IO/Q4UX9.
10. Diabetes Technology Network UK. DTN-UK statement on large language models. Association of British Clinical Diabetologists. January 2026. <https://abcd.care/dtn/resource/current/dtn-uk-statement-large-language-models>. Accessed 22 August 2026.
11. International Hypoglycaemia Study Group. Glucose concentrations of less than 3.0 mmol/L (54 mg/dL) should be reported in clinical trials: a joint position statement of the American Diabetes Association and the European Association for the Study of Diabetes. *Diabetologia.* 2017;60:3-6. <https://doi.org/10.1007/s00125-016-4146-6>
12. Yuan J, Li H, Ding X, et al. Understanding and mitigating numerical sources of nondeterminism in LLM inference. In: *Advances in Neural Information Processing Systems* 38 (NeurIPS 2025). 2025. <https://openreview.net/forum?id=Q3qAsZAEZw>

Tables

Table 1. Acceptance of the final response by the application, by model: conversations; accepted at first paste, with Wilson 95 per cent interval; accepted after the repair prompt; rejected with no prompt offered (invalid JSON or a required field absent); still rejected after repair; accepted responses breaking a rule the prompts state but the application does not check; stability index over accepted conversations.

Model	n	Accepted first paste, n (per cent; 95 per cent interval)	After repair	No prompt offered	Still rejected	Accepted but breaking a prompt rule	Stability index
Gemini 3.6 Flash	50	50 (100; 92.9 to 100)	0	0	0	4 of 50	0.951
Gemini 3.1 Pro	50	50 (100; 92.9 to 100)	0	0	0	3 of 50	0.983
Gemini 3.5 Flash-Lite	50	38 (76; 62.6 to 85.7)	1	11	0	6 of 39	0.989
GPT-5.6-sol	50	50 (100; 92.9 to 100)	0	0	0	0 of 50	0.983

Model	n	Accepted first paste, n (per cent; 95 per cent interval)	After repair	No prompt offered	Still rejected	Accepted but breaking a prompt rule	Stability index
GPT-5.4-mini	50	48 (96; 86.5 to 98.9)	0	1	1	25 of 48	0.998
Claude Opus 5	20	10 (50; 29.9 to 70.1)	1	9	0	5 of 11	0.972
Claude Sonnet 5	20	18 (90; 69.9 to 97.2)	1	1	0	2 of 19	0.998
Claude Haiku 4.5	50	17 (34; 22.4 to 47.8)	3	3	27	17 of 20	0.928
DeepSeek V4 Pro	50	45 (90; 78.6 to 95.7)	3	2	0	14 of 48	0.915
Grok 4.6	50	50 (100; 92.9 to 100)	0	0	0	1 of 50	0.986
Llama 4 Maverick	50	19 (38; 25.9 to 51.8)	25	2	4	5 of 44	0.953

Table 2. Settings any model changed in more than 10 per cent of accepted conversations. Each cell gives the percentage of accepted conversations proposing a change, with the range of suggested values in parentheses; 0 means no conversation proposed a change. Accepted conversations: Gemini 3.6 Flash 50, Gemini 3.1 Pro 50, Gemini 3.5 Flash-Lite 39, GPT-5.6-sol 50, GPT-5.4-mini 48, Claude Opus 5 11, Claude Sonnet 5 19, Claude Haiku 4.5 20, DeepSeek V4 Pro 48, Grok 4.6 50, Llama 4 Maverick 44. Current values: maximum IOB 25 U; maximum basal 12 U/h; target at 00:00 90 mg/dl; basal at 00:00 0.6 U/h; ISF 56 mg/dl per unit; dynamic ISF factor 70 per cent; DIA 10 h. AndroidAPS code defaults: maximum IOB 3 U, maximum basal 1 U/h, dynamic ISF factor 100 per cent; profile rows have no code default.

Model	Max IOB (U)	Max basal (U/h)	Target 00:00 (mg/dl)	Basal 00:00 (U/h)	ISF (mg/dl per unit)	Dynamic ISF factor (per cent)	DIA (h)
Gemini 3.6 Flash	98 (4 to 15)	96 (2.5 to 4)	78 (100)	22 (0.5)	0	26 (50)	0
Gemini 3.1 Pro	90 (3 to 10)	82 (2.5 to 4)	90 (100 to 110)	2 (0.5)	0	0	0
Gemini 3.5 Flash-Lite	13 (6 to 15)	0	3 (95)	10 (0.5)	0	0	0
GPT-5.6-sol	0	0	36 (100)	0	0	0	0
GPT-5.4-mini	0	0	0	4 (0.55)	0	0	0
Claude Opus 5	73 (6 to 10)	18 (4 to 5)	0	0	0	18 (55 to 60)	0
Claude Sonnet 5	0	0	0	0	0	0	0
Claude Haiku 4.5	5 (20)	0	0	15 (0.5 to 0.55)	10 (50 to 110)	0	5 (6)

Model	Max IOB (U)	Max basal (U/h)	Target 00:00 (mg/dl)	Basal 00:00 (U/h)	ISF (mg/dl per unit)	Dynamic ISF factor (per cent)	DIA (h)
DeepSeek V4 Pro	44 (2 to 12)	19 (2 to 4)	2 (100)	42 (0.5 to 0.55)	42 (30 to 180)	19 (0 to 100)	17 (5 to 8)
Grok 4.6	0	0	20 (100)	0	0	0	0
Llama 4 Maverick	39 (15 to 20)	2 (10)	0	52 (0.55 to 0.65)	9 (50 to 55)	0	0

Table 3. How the accepted conversations read the overnight lows: attribution to earlier insulin (the tail of evening dosing, stacking, carry-over or residual IOB) or to the target or basal itself, against whether the conversation changed the target at 00:00 or 23:00 or the basal at 00:00. A conversation may attribute to both. Wilson 95 per cent intervals in brackets.

Model	n	Attributes lows to earlier insulin	Attributes lows to target or basal	Changed overnight target or basal	Changed despite attributing to earlier insulin	Changed with no causal attribution
Gemini 3.6 Flash	50	21 (42; 29 to 56)	45 (90; 79 to 96)	44 (88; 76 to 94)	16	3
Gemini 3.1 Pro	50	40 (80; 67 to 89)	42 (84; 71 to 92)	45 (90; 79 to 96)	36	2
Gemini 3.5 Flash-Lite	39	1 (3; 1 to 13)	16 (41; 27 to 57)	5 (13; 6 to 27)	0	2
GPT-5.6-sol	50	35 (70; 56 to 81)	47 (94; 84 to 98)	18 (36; 24 to 50)	9	1
GPT-5.4-mini	48	14 (29; 18 to 43)	35 (73; 59 to 83)	2 (4; 1 to 14)	0	0
Claude Opus 5	11	11 (100; 74 to 100)	11 (100; 74 to 100)	0 (0; 0 to 26)	0	0
Claude Sonnet 5	19	18 (95; 75 to 99)	14 (74; 51 to 88)	0 (0; 0 to 17)	0	0
Claude Haiku 4.5	20	15 (75; 53 to 89)	19 (95; 76 to 99)	3 (15; 5 to 36)	1	0
DeepSeek V4 Pro	48	26 (54; 40 to 67)	47 (98; 89 to 100)	20 (42; 29 to 56)	6	1
Grok 4.6	50	34 (68; 54 to 79)	42 (84; 71 to 92)	10 (20; 11 to 33)	6	0
Llama 4 Maverick	44	0 (0; 0 to 8)	5 (11; 5 to 24)	23 (52; 38 to 66)	0	18

Figures




Figure 1

Figure 1. The application's workflow as replicated: the package and four prompts go to a new conversation in the user's chosen model; the fourth response is pasted back; the application accepts it, issues a repair prompt for three error types, or stops without a prompt when the paste is not valid JSON or a required field is absent.




Figure 2

Figure 2. Suggested values in accepted conversations that proposed a change, by model, for six settings, against the current value and the AndroidAPS code default; the count at the right of each row is conversations that changed the setting over conversations accepted.




Figure 3

Figure 3. Every conversation as one cell, in the order run. Top: the maximum IOB limit each conversation proposed, by model, with darker blue for lower values, grey where the limit was left at 25 U and hatching where the application rejected the response; the boxed cell marks one conversation, which is what a user receives. Middle: the ISF proposed by DeepSeek V4 Pro in the 20 conversations that changed it, against the current 56 mg/dl per unit. Bottom: the focus of Gemini 3.6 Flash's primary recommendation in each of its 50 conversations.

Supplementary material

Table S1. Model identifiers as requested and as returned by each provider

The requested identifier is the string sent in the API call; the served identifier is the model field the provider returned in each response, with the number of turns reporting it. Dates are the first and last conversation start dates (UTC).

Model (paper)	Provider and route	Requested identifier	Served identifier (turns)	Conversations	Dates
Gemini 3.6 Flash	Google, Generative Language API	gemini-3.6-flash	gemini-3.6-flash (204)	50	2026-08-21 to 2026-08-21
Gemini 3.1 Pro Preview	Google, Generative Language API	gemini-3.1-pro-preview	gemini-3.1-pro-preview (203)	50	2026-08-21 to 2026-08-21
Gemini 3.5 Flash-Lite	Google, Generative Language API	gemini-3.5-flash-lite	gemini-3.5-flash-lite (218)	50	2026-08-22 to 2026-08-22
GPT-5.6-sol	OpenRouter, OpenAI-compatible API (reasoning effort medium)	gpt-5.6-sol	gpt-5.6-sol (200)	50	2026-08-21 to 2026-08-21
GPT-5.4-mini	OpenRouter, OpenAI-compatible API (reasoning effort medium)	gpt-5.4-mini	gpt-5.4-mini-2026-03-17 (227)	50	2026-08-22 to 2026-08-22
Claude Opus 5	Anthropic, Messages API	claude-opus-5	claude-opus-5 (95)	20	2026-08-21 to 2026-08-22
Claude Sonnet 5	Anthropic, Messages API	claude-sonnet-5	claude-sonnet-5 (84)	20	2026-08-21 to 2026-08-22
Claude Haiku 4.5	Anthropic, Messages API	claude-haiku-4-5	claude-haiku-4-5-20251001 (244)	50	2026-08-22 to 2026-08-22
DeepSeek V4 Pro	OpenRouter, OpenAI-compatible API	deepseek/deepseek-v4-pro	deepseek/deepseek-v4-pro (217)	50	2026-08-22 to 2026-08-22
Grok 4.6	OpenRouter, OpenAI-compatible API	x-ai/grok-4.6	x-ai/grok-4.6 (201)	50	2026-08-22 to 2026-08-22
Llama 4 Maverick	OpenRouter, OpenAI-compatible API	meta-llama/llama-4-maverick	meta-llama/llama-4-maverick (231)	50	2026-08-22 to 2026-08-22

Table S2. Verification of the reconstructed validator against the application

Twenty-six real model outputs were pasted into the application’s result screen on 22 August 2026 and its response recorded; the same texts were then judged by the validator. Sample 04 was pasted while the application held the demonstration package and is reported only for whether fenced JSON parsed. The verdict classes are: accepted; repair prompt (with the clauses the application composed); schema error with no prompt; invalid JSON with no prompt.

Sample	Source conversation	Defect in the paste	Application's response	Validator's verdict	Agree
01	real_gemini35flashlite_50/run_037	fully compliant, should be accepted	accepted	accepted	yes
02	real_gpt54mini_med_50/run_018	current_value not copied verbatim: parameter_decisions.aaps.core.carb_absorption.current_value	repair prompt: copy current_value for aaps.core.carb_absorption	repair prompt: copy current_value for aaps.core.carb_absorption	yes
03	real_gpt54mini_med_50/run_008	unrequested key added inside meal_strategy_summary	accepted	accepted	yes
04	real_llama4maverick_50/run_035	fully compliant JSON wrapped in json fences accepted (fenced JSON parsed; tested against the demonstration package) parsed (fenced JSON accepted by extractor) yes 05 real_gemini35flashlite_50/run_026 confidence outside low/medium/high in: analysis_steps.meal_ecarbs_absorption.confidence accepted accepted yes 06 real_deepseek_v4pro_50/run_041 profile_recommendation.focus outside basal/target/cr/isf/dia repair prompt: profile_recommendation.foc	accepted	accepted	yes

Sample	Source conversation	Defect in the paste	Application's response	Validator's verdict	Agree
		us repair prompt: profile_recom mendation.foc us yes 07 real_opus5/ru n_017 recommendatio n priority outside low/medium/hi gh/safety accepted accepted yes 08 real_gemini31 pro_50/run_02 4 top-level section(s) missing: issues accepted accepted yes 09 real_gpt54min i_med_50/run_ 043 sections flagged: analysis_step s.automation_ sensitivity_l imits.confide nce, meal_strategy _summary accepted accepted yes 10 real_deepseek _v4pro_50/run _023 sections flagged: analysis_step s.automation_ smb_uam.confi dence, profile_recom mendation.foc us repair prompt: profile_recom mendation.foc			

Sample	Source conversation	Defect in the paste	Application's response	Validator's verdict	Agree
		us repair prompt: profile_recom mendation.foc us yes 11 real_gemini35 flashlite_50/ run_050 sections flagged: analysis_step s.basal_targe ts.confidence , analysis_step s.synthesis_p lan.confidenc e, profile_recom mendation schema error, no prompt (application session then crashed) schema error, no prompt yes 12 real_deepseek _v4pro_50/run _033 sections flagged: analysis_step s.profile_cr_ isf_dia.confid ence, parameter_dec isions.profil e.basal.00_00 .current_valu e, parameter_dec isions.profil e.basal.05_00 .current_valu e, parameter_dec isions.profil e.basal.06_00 .current_valu e, parameter_dec isions.profil			

Sample	Source conversation	Defect in the paste	Application's response	Validator's verdict	Agree
		e.basal.10_00 .current_valu e, parameter_dec isions.profil e.basal.16_00 .current_valu e, parameter_dec isions.profil e.basal.18_00 .current_valu e, parameter_dec isions.profil e.basal.19_00 .current_valu e, parameter_dec isions.profil e.cr.00_00.cu rrent_value, parameter_dec isions.profil e.cr.04_00.cu rrent_value, parameter_dec isions.profil e.cr.08_00.cu rrent_value, parameter_dec isions.profil e.cr.16_00.cu rrent_value, parameter_dec isions.profil e.cr.20_00.cu rrent_value, parameter_dec isions.profil e.isf.00_00.c urrent_value, parameter_dec isions.profil e.target.00_0 0.current_val ue, parameter_dec isions.profil e.target.08_0 0.current_val ue, parameter_dec isions.profil e.target.23_0			

Sample	Source conversation	Defect in the paste	Application's response	Validator's verdict	Agree
		0.current_value, parameter_decisions.profile.dia.current_value repair prompt: copy current_value for 17 keys repair prompt: copy current_value for 17 keys yes 13 real_grok46_50/run_017 sections flagged: analysis_steps.safety_overview, analysis_steps.basal_targets, analysis_steps.profile_crisf_dia, analysis_steps.automation_smb_uam, analysis_steps.automation_sensitivity_limits, analysis_steps.meal_bolus_strategy, analysis_steps.meal_ecarbs_absorption, analysis_steps.synthesis_plan accepted accepted yes 14 real_gpt54mini_med_50/run_003 sections flagged: meal_strategy_summary.confidence			

Sample	Source conversation	Defect in the paste	Application's response	Validator's verdict	Agree
		accepted accepted yes 15 real_gemini35 flashlite_50/ run_015 sections flagged: parameter_dec isions.aaps.c ore.carb_abso rption accepted accepted yes 16 real_deepseek _v4pro_50/run _002 sections flagged: parameter_dec isions.profil e.basal.10_00 , parameter_dec isions.profil e.basal.16_00 schema error, no prompt schema error, no prompt yes 17 real_deepseek _v4pro_50/run _035 sections flagged: parameter_dec isions.profil e.cr.00_00 accepted accepted yes 18 real_gemini36 flash_50/run_ 016 sections flagged: parameter_dec isions.profil e.dia accepted accepted yes 19			

Sample	Source conversation	Defect in the paste	Application's response	Validator's verdict	Agree
		real_gpt54mini_med_50/run_042 sections flagged: profile_recommendation accepted accepted yes 20 real_gemini35flashlite_50/run_017 valid JSON followed by a stray closing			
21	real_gemini35flashlite_50/run_028	well-formed except one closing brace missing (root never closes)	invalid JSON, no prompt	invalid JSON, no prompt	yes
22	real_opus5/run_011	several top-level JSON objects in sequence instead of one	invalid JSON, no prompt	invalid JSON, no prompt	yes
23	real_llama4maverick_50/run_012	syntax error inside: Expecting value at char 0	invalid JSON, no prompt	invalid JSON, no prompt	yes
24	real_gemini35flashlite_50/run_025	stray closing `` at the end, no opening fence	invalid JSON, no prompt	invalid JSON, no prompt	yes
25	real_gemini35flashlite_50/run_033	parameter_decisions rows missing/duplicated/extra: parameter_decisions	repair prompt: add profile.cr.16_00, remove profile.cr.06_00	repair prompt: add profile.cr.16_00; remove profile.cr.06_00	yes
26	real_gpt54mini_med_50/run_018 (response to the repair prompt)	current_value still not copied exactly ("5m" for "5 min")	repair prompt: copy current_value for aaps.core.carb_absorption (second time)	repair prompt: copy current_value for aaps.core.carb_absorption	yes

Agreement on verdict class: 25 of 25 informative samples.

Table S3. Whether the rationale for each decision cites a figure from the record

Post hoc analysis over accepted conversations. For each decision row the rationale text was searched for figures with a unit (the same extraction as the citation check) and each figure was checked against the numbers in the package. A change decision is counted as cited when at least

one such figure is present in the package. Wilson 95 per cent intervals, with the decision row as the unit, are given for the cited proportion; rows are not independent within a conversation, so the intervals are indicative.

Model	Change decisions	Cited a verified figure, n (per cent; 95 per cent interval)	Cited figures, none verified	No figure in rationale	Keep decisions	Keep decisions citing a verified figure, n (per cent)
Gemini 3.6 Flash	160	154 (96; 92 to 98)	0	6	990	328 (33)
Gemini 3.1 Pro Preview	141	97 (69; 61 to 76)	0	44	1009	88 (9)
Gemini 3.5 Flash-Lite	10	9 (90; 60 to 98)	0	1	887	129 (15)
GPT-5.6-sol	20	17 (85; 64 to 95)	0	3	1130	187 (17)
GPT-5.4-mini	2	0 (0; 0 to 66)	0	2	1102	3 (0)
Claude Opus 5	16	16 (100; 81 to 100)	0	0	237	200 (84)
Claude Sonnet 5	1	1 (100; 21 to 100)	0	0	436	150 (34)
Claude Haiku 4.5	27	27 (100; 88 to 100)	0	0	433	405 (94)
DeepSeek V4 Pro	93	84 (90; 83 to 95)	0	9	1011	105 (10)
Grok 4.6	16	16 (100; 81 to 100)	0	0	1134	615 (54)
Llama 4 Maverick	56	3 (5; 2 to 15)	0	53	956	0 (0)

Table S4. How accepted conversations read the package’s bolus-only insulin total

Post hoc analysis. The package’s treatment_summary gives total_insulin_u 275.95 over 14 days (19.7 U/day), a sum of the bolus and SMB treatment records only: the package carries no basal delivery data, and its scheduled basal profile integrates to 16.2 U/day. Response text was searched for the derived figure (19.7 or 275.95) in a daily total or TDD context, for statements that the total excludes basal or is bolus only, and for a reconstructed total near 36 U/day. Wilson 95 per cent intervals over accepted conversations.

Model	Accepted	Treated 19.7 U as a daily total or TDD, n (per cent; 95 per cent interval)	Noted the total excludes basal, n (per cent)	Reconstructed about 36 U/day, n (per cent)	Lowered max IOB while calling 19.7 a daily total
Gemini 3.6 Flash	50	49 (98; 90 to 100)	13 (26)	0 (0)	48

Model	Accepted	Treated 19.7 U as a daily total or TDD, n (per cent; 95 per cent interval)	Noted the total excludes basal, n (per cent)	Reconstructed about 36 U/day, n (per cent)	Lowered max IOB while calling 19.7 a daily total
Gemini 3.1 Pro Preview	50	16 (32; 21 to 46)	5 (10)	0 (0)	16
Gemini 3.5 Flash-Lite	39	0 (0; 0 to 9)	5 (13)	0 (0)	0
GPT-5.6-sol	50	4 (8; 3 to 19)	23 (46)	2 (4)	0
GPT-5.4-mini	48	0 (0; 0 to 7)	32 (67)	0 (0)	0
Claude Opus 5	11	9 (82; 52 to 95)	10 (91)	11 (100)	7
Claude Sonnet 5	19	2 (11; 3 to 31)	8 (42)	6 (32)	0
Claude Haiku 4.5	20	4 (20; 8 to 42)	19 (95)	8 (40)	0
DeepSeek V4 Pro	48	5 (10; 4 to 22)	25 (52)	7 (15)	5
Grok 4.6	50	26 (52; 38 to 65)	23 (46)	29 (58)	0
Llama 4 Maverick	44	0 (0; 0 to 8)	0 (0)	0 (0)	0